

Optimizers(Gradient Descent)

```
X = [0.5, 2.5]
Y = [0.2, 0.9]

def f(w,b,x) : #sigmoid with parameters w,b
    return 1.0 / (1.0 + np.exp(-(w*x + b)))



def error (w, b) :
    err = 0.0
    for x,y in zip(X,Y) :
        fx = f(w,b,x)
        err += 0.5 * (fx - y) ** 2
    return err

def grad_b(w,b,x,y) :
    fx = f(w,b,x)
    return (fx - y) * fx * (1 - fx)

def grad_w(w,b,x,y) :
    fx = f(w,b,x)
    return (fx - y) * fx * (1 - fx) * x


def do_gradient_descent() :
    w, b, eta, max_epochs = -2, -2, 1.0, 1000
    for i in range(max_epochs) :
        dw, db = 0, 0
        for x,y in zip(X, Y) :
            dw += grad_w(w, b, x, y)
            db += grad_b(w, b, x, y)
        w = w - eta * dw
        b = b - eta * db
```

Gradient descent on the error surface




Optimizers(Momentum Gradient Descent)

```
def do_momentum_gradient_descent() :  
    w, b, eta = init_w, init_b, 1.0  
    prev_v_w, prev_v_b, gamma = 0, 0, 0.9  
    for i in range(max_epochs) :  
        dw, db = 0, 0  
        for x,y in zip(X, Y) :  
            dw += grad_w(w, b, x, y)  
            db += grad_b(w, b, x, y)  
  
        v_w = gamma * prev_v_w + eta * dw  
        v_b = gamma * prev_v_b + eta * db  
        w = w - v_w  
        b = b - v_b  
        prev_v_w = v_w  
        prev_v_b = v_b
```




Optimizers(Nesterov Accelerated Gradient Descent)

```
def do_nesterov_accelerated_gradient_descent() :  
  
    w, b, eta = init_w, init_b , 1.0  
    prev_v_w, prev_v_b, gamma = 0, 0, 0.9  
    for i in range(max_epochs) :  
        dw, db = 0, 0  
        #do partial updates  
        v_w = gamma * prev_v_w  
        v_b = gamma * prev_v_b  
        for x,y in zip(X, Y) :  
            #calculate gradients after partial update  
            dw += grad_w(w - v_w, b - v_b, x, y)  
            db += grad_b(w - v_w, b - v_b, x, y)  
  
        #now do the full update  
        v_w = gamma * prev_v_w + eta * dw  
        v_b = gamma * prev_v_b + eta * db  
        w = w - v_w  
        b = b - v_b  
        prev_v_w = v_w  
        prev_v_b = v_b
```




Optimizers(Adagrad)

```
def do_adagrad():  
    w, b, eta = init_w, init_b, 0.1  
    v_w, v_b, eps = 0, 0, 1e-8  
    for i in range(max_epochs):  
        dw, db = 0, 0  
        for x,y in zip(X, Y):  
            dw += grad_w(w, b, x, y)  
            db += grad_b(w, b, x, y)  
  
        v_w = v_w + dw**2  
        v_b = v_b + db**2  
  
        w = w - (eta / np.sqrt(v_w + eps)) * dw  
        b = b - (eta / np.sqrt(v_b + eps)) * db
```




Optimizers(RmsProp)

```
def do_rmsprop() :  
    w, b, eta = init_w, init_b, 0.1  
    v_w, b_updates, eps, beta1 = 0, 0, 1e-8, 0.9  
    for i in range(max_epochs) :  
        dw, db = 0, 0  
        for x,y in zip(X, Y) :  
            dw += grad_w(w, b, x, y)  
            db += grad_b(w, b, x, y)  
  
        v_w = beta1 * v_w + (1 - beta1) dw**2  
        v_b = beta1 * v_b + (1 - beta1) db**2  
  
        w = w - (eta / np.sqrt(v_w + eps)) * dw  
        b = b - (eta / np.sqrt(v_b + eps)) * db
```



Optimizers(Adam)

```
def do_adam() :  
    w_b_dw_db = [(init_w, init_b, 0, 0)]  
    w_history, b_history, error_history = [], [], []  
    w, b, eta, mini_batch_size, num_points_seen =  
        init_w, init_b, 0.1, 10, 0  
    m_w, m_b, v_w, v_b, m_w_hat, m_b_hat, v_w_hat,  
        v_b_hat, eps, beta1, beta2 = 0, 0, 0, 0, 0,  
        0, 0, 0, 1e-8, 0.9, 0.999  
    for i in range(max_epochs) :  
        dw, db = 0, 0  
        for x,y in zip(X, Y) :  
            dw += grad_w(w, b, x, y)  
            db += grad_b(w, b, x, y)  
  
        m_w = beta1 * m_w + (1-beta1)*dw  
        m_b = beta1 * m_b + (1-beta1)*db  
  
        v_w = beta2 * v_w + (1-beta2)*dw**2  
        v_b = beta2 * v_b + (1-beta2)*db**2  
  
        m_w_hat = m_w/(1-math.pow(beta1,i+1))  
        m_b_hat = m_b/(1-math.pow(beta1,i+1))  
  
        v_w_hat = v_w/(1-math.pow(beta2,i+1))  
        v_b_hat = v_b/(1-math.pow(beta2,i+1))  
  
        w = w - (eta / np.sqrt(v_w_hat + eps)) *  
            m_w_hat  
        b = b - (eta / np.sqrt(v_b_hat + eps)) *  
            m_b_hat
```



Stochastic ,Batch ,Mini-Batch

```
def do_stochastic_gradient_descent():
    w, b, eta, max_epochs = -2, -2, 1.0, 1000
    for i in range(max_epochs):
        dw, db = 0, 0
        for x, y in zip(X, Y):
            dw = grad_w(w, b, x, y)
            db = grad_b(w, b, x, y)
            w = w - eta * dw
            b = b - eta * db
```

```
def do_gradient_descent() :
    w, b, eta, max_epochs = -2, -2, 1.0, 1000
    for i in range(max_epochs) :
        dw, db = 0, 0
        for x,y in zip(X, Y) :
            dw += grad_w(w, b, x, y)
            db += grad_b(w, b, x, y)
        w = w - eta * dw
        b = b - eta * db
```

```
def do_mini_batch_gradient_descent() :
    w, b, eta = -2, -2, 1.0
    mini_batch_size, num_points_seen = 2, 0
    for i in range(max_epochs) :
        dw, db, num_points = 0, 0, 0
        for x,y in zip(X, Y) :
            dw += grad_w(w, b, x, y)
            db += grad_b(w, b, x, y)
            num_points_seen += 1

        if num_points_seen % mini_batch_size == 0 :
            # seen one mini_batch
            w = w - eta * dw
            b = b - eta * db
            dw, db = 0, 0 #reset gradients
```